FILTER LABS

Ubiquity API

Detailed Reference

Requests, parameters and response structures from the OpenAPI 3.0.3 specification, v1.7.0

What is in this document

The detailed edition of the Ubiquity API guide: for all 56 operations, the parameters, request body fields, example payloads and response structures, grouped across 10 functional areas.

Overview

Ubiquity discovers open-source signals, geotags and verifies them, and turns them into live feeds you can monitor, query and route into existing workflows. The API covers the full lifecycle of that platform: discover, stream and understand.

56	10	46	JWT
operations	resource areas	unique paths	bearer auth

Servers

Base URL	Purpose
`https://ubiquity.filterlabs.ai/api`	Reverse-proxied backend (locations / orchestrator / understand services)
`https://auth.filterlabs.ai`	Keycloak auth server (realm filter-labs-web, client web-app)
`https://agents.ubiquity.filterlabs.ai`	Trial / preview signup site ('Ubiquity for Agents'): signup and health check only

Authentication

All endpoints except the token call and the two signup-site calls require a JWT, sent as `Authorization: Bearer ***`. Tokens come from the Keycloak token endpoint (realm filter-labs-web, client web-app) using the password grant, and are renewed with the refresh_token grant. Access tokens last about 30 minutes; refresh tokens last about 8 hours and rotate on use.

Service families

Prefix	Service	Handles
`/locations`	Locations	Pipelines, feeds, entities, jobs, saved searches, sharing, media, direct data stream.
`/orchestrator`	Orchestrator	Query parsing, discovery job launch, dataset-upload workflow.
`/understand`	Understand	Ubi chat, artifacts, metric previews.
`/api/v1/signup, /health`	Signup site	Trial account signup and health check (separate base URL, no token).

Conventions

Full replacements. Several update calls are full replacements, not partial patches. Feed `PUT`, metric generator `PUT` and pipeline `PATCH` (whole `metadata` object) all drop fields you omit, so read the resource first and resend unchanged values.

Identifiers. `feed_id`, `entity_id` and entity source ids are integers; `pipeline_id`, `saved_search_id`, `job_id` and `session_id` are strings. For the user id, entity votes take the Keycloak UUID (the JWT `sub` claim) while the chat endpoints and the permissions endpoints take the account email.

Paging. List endpoints take `limit` and `offset`. The feed stream is different: it pages with an opaque cursor (returned as `next_cursor`) and has no offset.

Status codes in use

Code	Meaning in this API
200 / 201	Success; 201 is returned when a resource is created (pipelines, routes, feeds, saved searches).
202	Accepted for background processing (source upload).
204	Success with an empty body (deletes of routes and metric generators).
400	Rejected request, for example `POST /pipelines` before its feed exists returns `{"error": "Feed not found"}`.
404	Not found, for example a media asset that is not stored, or an artifact that does not belong to the feed.
422	Validation error, returned as `{"detail": [{"loc": [...], "msg": "Field required"}]}`.
500	`POST /artifacts/preview` before any artifacts have been ingested (expected at that stage).

Core workflow: creating a pipeline

Pipeline creation is three ordered calls, each returning an id the next one needs. Launching a discovery job is a separate, optional fourth step.

Step	Method	Endpoint	What it does	Returns	Used by
1	**POST**	`/locations/api/v1/saved-searches`	Save the parsed query and filters.	id (saved_search_id)	POST /pipelines
2	**POST**	`/locations/api/v1/feeds`	Create the feed that will hold collected content.	id (feed_id)	POST /pipelines
3	**POST**	`/locations/v1/pipelines`	Link the saved search and feed into a new pipeline.	id (pipeline_id)	PATCH, routes, jobs/create
4	**POST**	`/orchestrator/api/v1/jobs/create`	Optional. Start discovery; poll with GET /locations/api/v1/jobs/{job_id}.	id (job_id)	GET /jobs/{job_id}

Where the inputs come from

The query is first parsed by `POST /search/unified`. Its `location.id` becomes the saved search's `search_filters.location_uuid`, its parsed locations, topics and entity types are flattened into `search_filters.entities`, and its `session_id` is passed to `POST /pipelines`.

Importing your own sources

Add URLs with `POST /feeds/{feed_id}/sources/upload` (returns 202), or upload a behavioral dataset with `POST /orchestrator/workflows/dataset-upload/trigger` and poll its status endpoint until it completes.

Ubi chat in one call

Send `{"message": "..."}` to the session messages endpoint. The call is synchronous, so the reply, including any `filter_update`, comes back in the same response and no polling is needed. The session is created implicitly if it does not exist yet.

Checking ingestion progress

Use `GET /understand/artifacts/count/{feed_id}` for a fast, inexpensive count of ingested artifacts. It is cheaper than reading `total_count` from `artifacts/preview`.

Getting a trial account

Send an email address (first and last name are optional) to `POST /api/v1/signup` on the signup site; no token is needed. The call is idempotent: re-submitting an existing email returns the same `user_id` with `created=false` and does not send another invite.

Sharing a pipeline

A pipeline can only be shared with members of your own organizations; arbitrary email invites are not supported. List your organizations, list the members of one, then grant access with a `PUT` to the pipeline's permissions.

Reading media

Stream items from platforms that capture media carry a `media_assets` list. Pass an asset's ids to the media endpoint to get a presigned download URL that expires after about 5 minutes.

Endpoint reference

Methods: **GET** **POST** **PUT** **PATCH** **DELETE**. Example request bodies use placeholder values in angle brackets; response structures show field names with their types, and `...` where the type is not documented yet.

01 01 Authentication

1 endpoint

Keycloak-issued JWTs, sent as a Bearer token on every other call.

#	Method	Path	Description
1.1	POST	`/realms/filter-labs-web/protocol/openid-connect/token`	Log in or refresh a token (Keycloak password and refresh_token grants).

1.1 POST `/realms/filter-labs-web/protocol/openid-connect/token`

Log in or refresh a token (Keycloak password and refresh_token grants).

Send the body form-encoded. Use `grant_type=password` with `username` and `password` to log in, or `grant_type=refresh_token` with `refresh_token` to renew.

Request body `application/x-www-form-urlencoded`

Field	Type	Req	Notes
`grant_type`	string		One of: `password`, `refresh_token`. One of: `password`, `refresh_token`.
`client_id`	string		
`username`	string		
`password`	string		
`refresh_token`	string		

Example

```
grant_type=password&client_id=web-app&username=<email>&password=<password>

grant_type=refresh_token&client_id=web-app&refresh_token=<refresh_token>
```

Responses

200 Tokens issued. The access token lasts about 30 minutes; the refresh token about 8 hours and rotates, so always store the newest one.

```
{
  "access_token": "<JWT, ~30 min lifetime>",
  "refresh_token": "<token, ~8 h lifetime, rotates on use>",
  ...
}
```

02 02 Trial Signup & Health

2 endpoints

Create a trial account and check service health. Both calls are served by the signup site (a separate base URL) and need no token.

#	Method	Path	Description
2.1	POST	`/api/v1/signup`	Create a trial / preview account. No token needed; served by the signup site.
2.2	GET	`/health`	Health check for the signup site's backend. No token needed.

2.1 POST `/api/v1/signup`

Create a trial / preview account. No token needed; served by the signup site.

Served by the signup site (`https://agents.ubiquity.filterlabs.ai`), not the main backend. No Bearer token is needed.

Idempotent: submitting an email that already has an account returns the same `user_id` with `created` and `invited` false. No second invite is sent and no duplicate is created.

No deliverability or domain check is performed, so a syntactically valid but undeliverable address is created the same as a real one.

Request body `application/json`

Field	Type	Req	Notes
`email`	string (email)	Yes	
`first_name`	string		Optional. Omit the field entirely when blank; do not send an empty string.
`last_name`	string		Optional. Same omit-if-blank behavior as `first_name`.

Example

```
{
  "email": "<email>",
  "first_name": "<first_name>",
  "last_name": "<last_name>"
}
```

Responses

200 Account created, or already existing. `created` and `invited` are both true only for a genuinely new signup.

```
{
  "user_id": ...,
  "email": "<string>",
  "created": <boolean>,
  "licensed": ...,
  "invited": <boolean>
}
```

2.2 GET /health

Health check for the signup site's backend. No token needed.

Served by the signup site (`https://agents.ubiquity.filterlabs.ai`), not the main backend. No Bearer token is needed.

The signup API is hosted as part of the Scout search service, which is the service name reported here.

Responses

200 Service status.

```
{
  "status": ...,
  "service": "ubiquity-scout",
  "auth_disabled": ...,
  "twitter": { "configured": ... },
  "license": { "role": ..., "enforced": ... },
  "signup": { "configured": ... },
  "usage": { "daily_limit": ..., "backend": ... }
}
```

03 03 Search & Locations

3 endpoints

Turn a free-text query into structured entities and look up places.

#	Method	Path	Description
3.1	POST	`/orchestrator/search/unified`	Parse a free-text query into location, topic and entity-type entities. Also the refine loop for the Discover chat.
3.2	GET	`/locations/v1/locations`	Search locations already in Ubiquity by name (not a geocoder). Pages with limit and offset.
3.3	GET	`/locations/v1/locations/{location_uuid}`	Get a location object (may be a synthesized COMBINED-location).

3.1 POST /orchestrator/search/unified

Parse a free-text query into location, topic and entity-type entities. Also the refine loop for the Discover chat.

For the first turn send only `query` (plus optional limits). For each refine turn (change location, add a place, adjust topics and similar), call this same endpoint again and add `prior_query`, `prior_entities`, `prior_context` and `session_id`, each copied verbatim from the previous response.

Each parsed entity carries `value`, `type`, `confidence`, `normalized_value`, `aliases`, `metadata` and `languages`.

Request body `application/json`

Field	Type	Req	Notes
`query`	string		
`limit`	integer		
`max_new_sources`	integer		
`include_suggestions`	boolean		
`prior_query`	string		Refine turns only: previous turn's `query` string, verbatim.
`prior_entities`	object		Refine turns only: previous turn's `parsed_entities` object, verbatim.
`prior_context`	string		Refine turns only: previous turn's `context_report` string, verbatim.
`session_id`	string		Refine turns only: previous turn's response `session_id`, verbatim.

Example

```
{
  "query": "<query>",
  "limit": 20,
  "max_new_sources": 100,
  "include_suggestions": true,
  "prior_query": "<prior_query>",
  "prior_entities": {},
  "prior_context": "<prior_context>",
  "session_id": "<session_id>"
}
```

Responses

200 Parsed query. Pass `session_id` on to `POST /pipelines` and on to the next refine turn.

```
{
  "query": "<string>",
  "session_id": "<string>",
  "parsed_entities": {
    "locations": [...],
    "topics": [...],
    "entity_types": [...],
    "persons": [...],
    "organizations": [...],
    "events": [...]
  },
  "context_report": "<string>",
  "location": {
    "id": "<uuid>",
    "name": "<string>",
    "country": "<string>",
    "coordinates": ...,
    "bounding_box": ...,
    "is_multi_location": <boolean>,
    "component_locations": [...]
  },
  "entity_changes": { "added": [...], "removed": [...], "modified": [...] },
  "existing_sources": [...],
  "discovery_job_id": "<string>",
  "suggestions": [...],
  "stats": {...},
  "confidence": ...,
  "reasoning": "<string>"
}
```

3.2 GET `/locations/v1/locations`

Search locations already in Ubiquity by name (not a geocoder). Pages with limit and offset.

Searches only locations that already exist in Ubiquity. It is not a geocoder, and a place that has not been added through the unified search / discovery workflow will not appear.

`q` filters by name, and `offset` pages through results with no overlap between pages.

Parameters

Name	In	Type	Req	Notes
`q`	query	string	Yes	Search text.
`limit`	query	integer		Page size for list results.
`offset`	query	integer		Number of results to skip, for pagination.

Responses

200 Location matches, returned inside a double-nested envelope.

```
{
  "data": {
    "data": [ ...location matches... ]
  }
}
```

3.3 GET `/locations/v1/locations/{location_uuid}`

Get a location object (may be a synthesized COMBINED- location).

Parameters

Name	In	Type	Req	Notes
`location_uuid`	path	string	Yes	Location UUID.

Responses

200 The location.

```
{
  "id": "<string>",
  "name": "<string>",
  "abbreviation": "<string>",
  "country": "<string>",
  "coordinates": ...,
  "bounding_box": ...
}
```

04 04 Pipelines & Routing

8 endpoints

Pipelines tie a saved search to a feed; routes pass results between pipelines.

#	Method	Path	Description
4.1	GET	`/locations/v1/pipelines`	List pipelines.
4.2	POST	`/locations/v1/pipelines`	Create a pipeline from an existing saved search and feed (step 3 of creation).
4.3	GET	`/locations/v1/pipelines/{pipeline_id}`	Get a pipeline: feed_id, agent_config, discovery_config, search_filters.
4.4	PUT	`/locations/v1/pipelines/{pipeline_id}`	Full update, for example a rename.
4.5	PATCH	`/locations/v1/pipelines/{pipeline_id}`	Update metadata.agent_config and discovery_config. This is the Agent Settings save call; resend the whole metadata object.
4.6	GET	`/locations/v1/pipelines/{pipeline_id}/routes`	List routes (fan-in / fan-out) for a pipeline.
4.7	POST	`/locations/v1/pipelines/{pipeline_id}/routes`	Create a route from this pipeline to a destination pipeline.
4.8	DELETE	`/locations/v1/pipelines/{pipeline_id}/routes/{route_id}`	Delete a route.

4.1 GET /locations/v1/pipelines

List pipelines.

Parameters

Name	In	Type	Req	Notes
`limit`	query	integer		Page size for list results.
`offset`	query	integer		Number of results to skip, for pagination.
`order_by`	query	string		Field to sort by, for example `created_at`.
`order`	query	string		One of: `ASC`, `DESC`. One of: `ASC`, `DESC`.
`saved_search_id`	query	string		Saved-search id (returned by `POST /saved-searches`).

Responses

200 Array of pipeline objects. Shape of one element:

```
{
  "id": "<string>",
  "feed_id": <integer>,
  "saved_search_id": "<string>",
  "session_id": "<string>",
  "auto_sync_enabled": <boolean>,
  "sync_on_job_completion": <boolean>,
  "sync_on_map_change": <boolean>,
  "metadata": {},
  "created_at": "<timestamp>",
  "updated_at": "<timestamp>"
}
```

4.2 POST /locations/v1/pipelines

Create a pipeline from an existing saved search and feed (step 3 of creation).

Must follow `POST /saved-searches` and `POST /feeds`. Calling it alone fails with a 400 `{"error": "Feed not found"}`.

Creating a pipeline does not start discovery. Launching discovery is a separate, optional step (`PATCH` the pipeline, then `POST /jobs/create`).

Request body `application/json`

Field	Type	Req	Notes
`name`	string		
`saved_search_id`	string		Id from the preceding saved-searches POST.
`feed_id`	integer		Id from the preceding feeds POST.
`auto_sync_enabled`	boolean		Default: `true`.
`sync_on_job_completion`	boolean		Default: `true`.
`sync_on_map_change`	boolean		Default: `false`.
`session_id`	string		`session_id` from the originating `/search/unified` response.

Example

```
{
  "name": "<name>",
  "saved_search_id": "<saved_search_id>",
  "feed_id": 0,
  "auto_sync_enabled": true,
  "sync_on_job_completion": true,
  "sync_on_map_change": false,
  "session_id": "<session_id>"
}
```

Responses

201 Pipeline created.

```
{
  "id": "<string>",
  "feed_id": <integer>,
  "saved_search_id": "<string>",
  "session_id": "<string>",
  "auto_sync_enabled": <boolean>,
  "sync_on_job_completion": <boolean>,
  "sync_on_map_change": <boolean>,
  "metadata": {},
  "created_at": "<timestamp>",
  "updated_at": "<timestamp>"
}
```

4.3 GET `/locations/v1/pipelines/{pipeline_id}`

Get a pipeline: feed_id, agent_config, discovery_config, search_filters.

Parameters

Name	In	Type	Req	Notes
`pipeline_id`	path	string	Yes	Pipeline id.

Responses

200 The pipeline object. Once settings are saved, `metadata` carries `agent_config` and `discovery_config`; the object also exposes `feed_id` and `search_filters`.

```
{
  "id": "<string>",
  "feed_id": <integer>,
  "saved_search_id": "<string>",
  "session_id": "<string>",
  "auto_sync_enabled": <boolean>,
  "sync_on_job_completion": <boolean>,
  "sync_on_map_change": <boolean>,
  "metadata": {},
  "created_at": "<timestamp>",
  "updated_at": "<timestamp>"
}
```

4.4 PUT /locations/v1/pipelines/{pipeline_id}

Full update, for example a rename.

Parameters

Name	In	Type	Req	Notes
pipeline_id	path	string	Yes	Pipeline id.

Responses

200 The updated pipeline object.

```
{
  "id": "<string>",
  "feed_id": <integer>,
  "saved_search_id": "<string>",
  "session_id": "<string>",
  "auto_sync_enabled": <boolean>,
  "sync_on_job_completion": <boolean>,
  "sync_on_map_change": <boolean>,
  "metadata": {},
  "created_at": "<timestamp>",
  "updated_at": "<timestamp>"
}
```

4.5 PATCH /locations/v1/pipelines/{pipeline_id}

Update metadata.agent_config and discovery_config. This is the Agent Settings save call; resend the whole metadata object.

The whole `metadata` object must be re-sent; a partial diff is not supported.

Quick Presets (Balanced, Social Media Boost, News Focus, Behavioral Datasets Boost) are client-side form state only. Nothing is stored until this call is made.

Also used right before `POST /jobs/create` to set `discovery_config` defaults.

Parameters

Name	In	Type	Req	Notes
pipeline_id	path	string	Yes	Pipeline id.

Request body `application/json`

Field	Type	Req	Notes
`metadata`	object		
`metadata.agent_config`	object		
`metadata.agent_config.search_generation`	object		
`metadata.agent_config.search_generation.max_queries`	integer		Range 5 to 50.
`metadata.agent_config.search_generation.creativity_level`	string		One of: `low`, `medium`, `high`.
`metadata.agent_config.search_generation.focus_types`	array of string		
`metadata.agent_config.search_generation.language_preference`	string		One of: `native_first`, `mixed`, `english_first`.
`metadata.agent_config.search_generation.custom_instructions`	string		Free-text instructions for the search generation agent.
`metadata.agent_config.evaluation`	object		
`metadata.agent_config.evaluation.credibility_threshold`	number		Range 0 to 1.
`metadata.agent_config.evaluation.local_focus_priority`	string		One of: `low`, `medium`, `high`.
`metadata.agent_config.evaluation.quality_standards`	string		One of: `lenient`, `moderate`, `strict`.
`metadata.agent_config.evaluation.topic_relevance_weight`	number		Range 0 to 1.
`metadata.agent_config.datasets_discovery`	object		
`metadata.agent_config.datasets_discovery.enable_dataset_hunting`	boolean		
`metadata.agent_config.datasets_discovery.providers`	array of string		
`metadata.discovery_config`	object		
`metadata.discovery_config.max_queries`	integer		
`metadata.discovery_config.max_sources`	integer		
`metadata.discovery_config.min_sources`	integer		

Example

```
{
  "metadata": {
    "agent_config": {
      "search_generation": {
        "max_queries": 0,
        "creativity_level": "low",
        "focus_types": [
          "news"
        ],
        "language_preference": "native_first",
        "custom_instructions": "<custom_instructions>"
      },
      "evaluation": {
        "credibility_threshold": 0.5,
        "local_focus_priority": "low",
        "quality_standards": "lenient",
        "topic_relevance_weight": 0.5
      },
      "datasets_discovery": {
        "enable_dataset_hunting": true,
        "providers": [
          "fred"
        ]
      }
    },
    "discovery_config": {
      "max_queries": 0,
      "max_sources": 0,
      "min_sources": 0
    }
  }
}
```

Responses

200 The full updated pipeline object, with the new metadata.

```
{
  "id": "<string>",
  "feed_id": <integer>,
  "saved_search_id": "<string>",
  "session_id": "<string>",
  "auto_sync_enabled": <boolean>,
  "sync_on_job_completion": <boolean>,
  "sync_on_map_change": <boolean>,
  "metadata": {},
  "created_at": "<timestamp>",
  "updated_at": "<timestamp>"
}
```

4.6 GET `/locations/v1/pipelines/{pipeline_id}/routes`

List routes (fan-in / fan-out) for a pipeline.

Parameters

Name	In	Type	Req	Notes
`pipeline_id`	path	string	Yes	Pipeline id.

Responses

200 Array of routes. Source and destination pipeline names and feed_id are denormalized onto each one.

```
{
  "id": "<string>",
  "source_pipeline_id": "<string>",
  "destination_pipeline_id": "<string>",
  "source_pipeline_name": "<string>",
  "destination_pipeline_name": "<string>",
  "feed_id": <integer>,
  "enabled": <boolean>,
  "filter_mode": "<string>",
  "filter_config": {}
}
```

4.7 POST `/locations/v1/pipelines/{pipeline_id}/routes`

Create a route from this pipeline to a destination pipeline.

`filter_mode` and `filter_config` are accepted and stored. Their effect on what flows through the route is not documented yet.

Parameters

Name	In	Type	Req	Notes
`pipeline_id`	path	string	Yes	Pipeline id.

Request body `application/json`

Field	Type	Req	Notes
`destination_pipeline_id`	string		
`filter_mode`	string		Defaults to `all`. Accepted and stored (for example `liked_only`). Effect on routing is not documented yet. Default: `all`.
`filter_config`	object		Defaults to `{}`. Accepted and stored. Effect on routing is not documented yet. Default: `{}`.

Example

```
{
  "destination_pipeline_id": "<destination_pipeline_id>",
  "filter_mode": "all",
  "filter_config": {}
}
```

Responses

201 Route created.

```
{
  "id": "<string>",
  "source_pipeline_id": "<string>",
  "destination_pipeline_id": "<string>",
  "enabled": <boolean>,
  "filter_mode": "<string>",
  "filter_config": {}
}
```

4.8 DELETE `/locations/v1/pipelines/{pipeline_id}/routes/{route_id}`

Delete a route.

Parameters

Name	In	Type	Req	Notes
`pipeline_id`	path	string	Yes	Pipeline id.
`route_id`	path	string	Yes	Route id.

Responses

204 Empty body.

05 05 Sharing & Permissions

4 endpoints

Find who a pipeline can be shared with, and manage who has access.

#	Method	Path	Description
5.1	GET	`/locations/v1/organizations`	List the organizations the current user belongs to.
5.2	GET	`/locations/v1/permissions/organization/{org_uuid}`	List members of an organization: the only valid pipeline-sharing targets.
5.3	GET	`/locations/v1/permissions/pipeline/{pipeline_id}`	List current sharing grants for a pipeline.
5.4	PUT	`/locations/v1/permissions/pipeline/{pipeline_id}/{email}`	Grant an organization member access to a pipeline.

5.1 GET `/locations/v1/organizations`

List the organizations the current user belongs to.

Sharing targets must already be members of one of these organizations; arbitrary email invites are not supported.

Responses

200 Array of organizations. Most accounts see exactly one. Shape of one element:

```
{
  "name": "<string>",
  "uuid": "<string>",
  "licensed_seats": <integer>
}
```

5.2 GET `/locations/v1/permissions/organization/{org_uuid}`

List members of an organization: the only valid pipeline-sharing targets.

Despite its name, `keycloak_user_id` here is the member's email, not a UUID. `ability` is the organization-level baseline, not access to a particular pipeline.

Parameters

Name	In	Type	Req	Notes
`org_uuid`	path	string	Yes	Organization UUID (from `GET /organizations`).

Responses

200 Array of organization members. Shape of one element:

```
{
  "ability": "<string>",
  "keycloak_user_id": "<the member's email>"
}
```

5.3 GET /locations/v1/permissions/pipeline/{pipeline_id}

List current sharing grants for a pipeline.

Same shape as the organization members response.

Parameters

Name	In	Type	Req	Notes
`pipeline_id`	path	string	Yes	Pipeline id.

Responses

200 Array of current sharing grants, in the same shape as the organization members response:

```
{
  "ability": "<string>",
  "keycloak_user_id": "<the member's email>"
}
```

5.4 PUT /locations/v1/permissions/pipeline/{pipeline_id}/{email}

Grant an organization member access to a pipeline.

Send the email as a plain path segment, not URL-encoded.

The Share dialog grants `write` when Edit is checked, and `write` implies read. The value for view-only access is not documented yet; try `read` and check the result with `GET /permissions/pipeline/{pipeline_id}`.

Parameters

Name	In	Type	Req	Notes
`pipeline_id`	path	string	Yes	Pipeline id.
`email`	path	string	Yes	Target user's email, sent as a plain (not URL-encoded) path segment.

Request body `application/json`

Field	Type	Req	Notes
`ability`	string		Use `write` (implies read). The value for view-only access is not documented yet.

Example

```
{
  "ability": "write"
}
```

Responses

200 Access granted.

06 06 Saved Searches & Feeds

11 endpoints

Feeds hold the stream of collected content and its media; metric generators score it.

#	Method	Path	Description
6.1	POST	`/locations/api/v1/feeds`	Create the feed resource (step 2 of pipeline creation).
6.2	GET	`/locations/api/v1/feeds/{feed_id}`	Get a feed, including its refresh interval.
6.3	PUT	`/locations/api/v1/feeds/{feed_id}`	Full update: rename, refresh and detection intervals, review settings. GET first and resend unchanged fields.
6.4	GET	`/locations/api/v1/feeds/{feed_id}/metric_prompts`	List metric generators for a feed.
6.5	POST	`/locations/api/v1/feeds/{feed_id}/metric_prompts`	Create a metric generator.
6.6	PUT	`/locations/api/v1/feeds/{feed_id}/metric_prompts/{metric_id}`	Edit a metric generator (full PUT only).
6.7	DELETE	`/locations/api/v1/feeds/{feed_id}/metric_prompts/{metric_id}`	Delete a metric generator.
6.8	GET	`/locations/api/v1/feeds/{feed_id}/stream`	Direct Data Feed API: raw content items for the full corpus, with date, text, entity and URL filters, plus media asset metadata.
6.9	GET	`/locations/api/v1/feeds/{feed_id}/media/{artifact_uuid}/{asset_id}`	Resolve a stored media asset from a stream item to a short-lived presigned download URL.
6.10	POST	`/locations/api/v1/saved-searches`	Persist the parsed query and filters (step 1 of pipeline creation).
6.11	PUT	`/locations/api/v1/saved-searches/{saved_search_id}`	Re-save a pipeline's search query and filters.

6.1 POST `/locations/api/v1/feeds`

Create the feed resource (step 2 of pipeline creation).

Must run before `POST /pipelines`. The returned id is the numeric `feed_id` used throughout the API.

Request body `application/json`

Field	Type	Req	Notes
`name`	string		
`metadata`	object		
`metadata.description`	string		Free-text description, for example `Feed for search: <query text>`.
`metadata.search_query`	string		

Example

```
{
  "name": "<name>",
  "metadata": {
    "description": "Feed for search: <query text>",
    "search_query": "<search_query>"
  }
}
```

Responses

201 Feed created.

```
{
  "id": <integer>,
  "uuid": "<string>",
  "name": "<string>",
  "refresh_interval": { "Microseconds": ..., "Days": ..., "Months": ..., "Valid": false },
  "metadata": "<base64>",
  "keycloak_user_id": "<string>",
  "created_at": "<timestamp>",
  "updated_at": "<timestamp>"
}
```

6.2 GET `/locations/api/v1/feeds/{feed_id}`

Get a feed, including its refresh interval.

Parameters

Name	In	Type	Req	Notes
`feed_id`	path	integer	Yes	Numeric feed id (returned by `POST /feeds`).

Responses

200 The feed. Here the refresh interval is returned as a human-readable string.

```
{
  "id": <integer>,
  "uuid": "<string>",
  "name": "<string>",
  "refresh_interval": "<human-readable, e.g. '12 hours'>",
  ...
}
```

6.3 PUT `/locations/api/v1/feeds/{feed_id}`

Full update: rename, refresh and detection intervals, review settings. GET first and resend unchanged fields.

Full PUT, not a partial update: omitted fields (including the feed's own `metadata`) are reset to null or default. GET the feed first and resend unchanged values.

`refresh_interval_minutes` is the same field behind the Enable Ubi toggle.

Parameters

Name	In	Type	Req	Notes
`feed_id`	path	integer	Yes	Numeric feed id (returned by `POST /feeds`).

Request body `application/json`

Field	Type	Req	Notes
`name`	string		Must be re-sent (full PUT).
`refresh_interval_minutes`	integer		0 = do not update; 720 = 12 h; 1440 = daily; 2880 = every 2 days; 10080 = weekly; about 43200 = monthly.
`deletion_detection_interval_minutes`	integer		Settable by direct PUT. Effect on background jobs is not documented yet.
`engagement_detection_interval_minutes`	integer		Settable by direct PUT. Effect on background jobs is not documented yet.
`automatic_entity_rediscovery_interval_minutes`	integer		Settable by direct PUT. Effect on background jobs is not documented yet.
`force_review_for_auto_discover`	boolean		Settable by direct PUT.
`backdate_limit`	integer		int32. Probably a day count (unit not documented yet). A date string such as `YYYY-MM-DD` is rejected with a 400.

Example

```
{
  "name": "<name>",
  "refresh_interval_minutes": 0,
  "deletion_detection_interval_minutes": 0,
  "engagement_detection_interval_minutes": 0,
  "automatic_entity_rediscovery_interval_minutes": 0,
  "force_review_for_auto_discover": true,
  "backdate_limit": 0
}
```

Responses

200 The updated feed. Interval fields come back as a struct.

```
{
  "id": <integer>,
  "name": "<string>",
  "refresh_interval": { "Microseconds": ..., "Days": ..., "Months": ..., "Valid": <boolean> },
  "refreshed_at": "<timestamp>",
  ...
}
```

6.4 GET `/locations/api/v1/feeds/{feed_id}/metric_prompts`

List metric generators for a feed.

Parameters

Name	In	Type	Req	Notes
`feed_id`	path	integer	Yes	Numeric feed id (returned by `POST /feeds`).

Responses

200 Array of metric generators, or null when the feed has none. Shape of one element:

```
{
  "id": <string>,
  "feed_id": <integer>,
  "name": "<string>",
  "metric_name": "<string>",
  "prompt": "<string>",
  "filter_prompt": "<string> | null",
  "filter_distance": <number> | null,
  "active": <boolean>,
  "created_at": "<timestamp>",
  "updated_at": "<timestamp>"
}
```

6.5 POST `/locations/api/v1/feeds/{feed_id}/metric_prompts`

Create a metric generator.

`filter_prompt` and `filter_distance` are accepted and stored. Their effect is not documented yet.

Parameters

Name	In	Type	Req	Notes
`feed_id`	path	integer	Yes	Numeric feed id (returned by `POST /feeds`).

Request body `application/json`

Field	Type	Req	Notes
`feed_id`	integer		Numeric feed id (returned by `POST /feeds`).
`name`	string		
`metric_name`	string		Slug of `name`, generated by the client.
`prompt`	string		
`active`	boolean		
`filter_prompt`	string		Intended as a semantic pre-filter that gates which artifacts get scored. Accepted and stored. Effect is not documented yet. May be null. May be null.
`filter_distance`	number		Likely a vector-distance cutoff for that pre-filter. Accepted and stored. Effect is not documented yet. May be null. May be null.

Example

```
{
  "feed_id": 0,
  "name": "<name>",
  "metric_name": "<metric_name>",
  "prompt": "<prompt>",
  "active": true,
  "filter_prompt": "<filter_prompt>",
  "filter_distance": 0.5
}
```

Responses

200 Metric generator created.

```
{
  "id": <string>,
  "feed_id": <integer>,
  "name": "<string>",
  "metric_name": "<string>",
  "prompt": "<string>",
  "filter_prompt": "<string> | null",
  "filter_distance": <number> | null,
  "active": <boolean>,
  "created_at": "<timestamp>",
  "updated_at": "<timestamp>"
}
```

6.6 PUT `/locations/api/v1/feeds/{feed_id}/metric_prompts/{metric_id}`

Edit a metric generator (full PUT only).

This resource takes full PUT only; there is no partial PATCH. Prefer this call over creating a replacement.

Parameters

Name	In	Type	Req	Notes
`feed_id`	path	integer	Yes	Numeric feed id (returned by `POST /feeds`).
`metric_id`	path	string	Yes	Metric generator id.

Responses

200 The updated metric generator.

```
{
  "id": <string>,
  "feed_id": <integer>,
  "name": "<string>",
  "metric_name": "<string>",
  "prompt": "<string>",
  "filter_prompt": "<string> | null",
  "filter_distance": <number> | null,
  "active": <boolean>,
  "created_at": "<timestamp>",
  "updated_at": "<timestamp>"
}
```

6.7 DELETE `/locations/api/v1/feeds/{feed_id}/metric_prompts/{metric_id}`

Delete a metric generator.

Parameters

Name	In	Type	Req	Notes
`feed_id`	path	integer	Yes	Numeric feed id (returned by `POST /feeds`).
`metric_id`	path	string	Yes	Metric generator id.

Responses

204 Empty body.

6.8 GET /locations/api/v1/feeds/{feed_id}/stream

Direct Data Feed API: raw content items for the full corpus, with date, text, entity and URL filters, plus media asset metadata.

Returns the full corpus with no 1000-item display cap, so it suits direct analysis of feed content.

Page through the corpus with `cursor`: pass `pagination.next_cursor` from each response back verbatim as `cursor` until `has_more` is false.

`media_assets` appears only on items whose source platform supports media capture; otherwise the key is absent (not an empty array). Use `GET /feeds/{feed_id}/media/{artifact_uuid}/{asset_id}` to download a stored asset.

Parameters

Name	In	Type	Req	Notes
feed_id	path	integer	Yes	feed_id.
limit	query	integer		Range 1 to 1000. Default: 100. Range 1 to 1000. Default: 100.
cursor	query	string		Opaque token (base64-encoded JSON) returned as next_cursor in the previous response. Pass it back verbatim; do not decode or construct it.
order_by	query	string		published_at can be null (no usable date, or unparseable), and items with a null published_at are silently skipped under the default ordering. scraped_at is set at ingest and is never null. Use order_by=scraped_at to retrieve every artifact in the feed; the default ordering does not guarantee completeness. One of: published_at, scraped_at. Default: published_at. One of: published_at, scraped_at. Default: published_at.
content_max_length	query	integer		Range 0 to 10000. Default: 500. Range 0 to 10000. Default: 500.
include_vectors	query	string		
date_from	query	string		Start date, exclusive. RFC3339 or YYYY-MM-DD.
date_to	query	string		End date, inclusive. RFC3339 or YYYY-MM-DD.
text_search	query	string		
entity_ids	query	string		
entity_types	query	string		Comma-separated: institution, individual.
urls	query	string		

Responses

200 A page of raw content items.

```
{
  "items": [
    {
      "uuid": "<string>",
      "headline": "<string>",
      "content": "<string, truncated to content_max_length>",
      "url": "<string>",
      "published_at": "<timestamp>",
      "scraped_at": "<timestamp>",
      "entity_id": <integer>,
      "entity_type": "<institution | individual>",
      "authors": [...],
      "metrics": {...},
      "media_assets": [
```

```
        {
          "asset_id": "<string>",
          "position": <integer>,
          "kind": "<image | video | thumbnail>",
          "status": "<stored | failed | unavailable_logged_out>",
          "source_url": "<string>",
          "storage_key": "<string>",
          "sha256": "<string>",
          "mime": "<string>",
          "bytes": <integer>,
          "duration_s": <number>
        }
      ]
    }
  ],
  "pagination": {
    "limit": <integer>,
    "total": <integer>,
    "has_more": <boolean>,
    "next_cursor": "<opaque string> | null"
  }
}
```

6.9 GET `/locations/api/v1/feeds/{feed_id}/media/{artifact_uuid}/{asset_id}`

Resolve a stored media asset from a stream item to a short-lived presigned download URL.

Only assets with status `stored` can be retrieved. Find them in the `media_assets` list of an item from `GET /feeds/{feed_id}/stream`.

`artifact_uuid` is the stream item's own `uuid`, not its `entity_id`.

The returned URL expires after about 5 minutes; request a new one if needed.

Parameters

Name	In	Type	Req	Notes
`feed_id`	path	integer	Yes	Numeric feed id. It must be the feed the artifact belongs to: a different feed returns 404 "Artifact not found", and an unknown feed returns 400 "Feed not found".
`artifact_uuid`	path	string	Yes	The stream item's own `uuid` field (not `entity_id`).
`asset_id`	path	string	Yes	The `media_assets` entry's `asset_id`.

Responses

200 Download link for the asset. The `url` is a presigned GET URL that needs no further auth and serves the raw file bytes directly.

```
{
  "asset_id": "<string>",
  "kind": "<string>",
  "mime": "<string>",
  "bytes": <integer>,
  "url": "<presigned URL, ~5 min lifetime>",
  "expires_at": "<timestamp>"
}
```

404 Returned as `{"error": "Media not found"}` when the asset's status is `failed` or `unavailable_logged_out` (only `stored` assets can be served), or `{"error": "Artifact not found"}` when the `artifact_uuid` does not belong to the given feed.

400 Returned as `{"error": "Feed not found"}` when the `feed_id` does not exist.

6.10 POST /locations/api/v1/saved-searches

Persist the parsed query and filters (step 1 of pipeline creation).

Request body `application/json`

Field	Type	Req	Notes
`name`	string		Query text, used as the default pipeline name.
`search_query`	string		Same query text.
`search_filters`	object		
`search_filters.entities`	array of object		Flattened `parsed_entities.locations` + `.topics` + `.entity_types` from the preceding `/search/unified` response, each entry's value/type/confidence/normalized_value/aliases/metadata/languages kept verbatim.
`search_filters.location_uuid`	string		`location.id` from `/search/unified`.

Example

```
{
  "name": "<name>",
  "search_query": "<search_query>",
  "search_filters": {
    "entities": [
      {}
    ],
    "location_uuid": "<location_uuid>"
  }
}
```

Responses

201 Saved search created.

```
{
  "id": "<string>",
  "name": "<string>",
  "search_query": "<string>",
  "search_filters": { ... },
  "is_default": <boolean>,
  "created_at": "<timestamp>",
  "updated_at": "<timestamp>"
}
```

6.11 PUT /locations/api/v1/saved-searches/{saved_search_id}

Re-save a pipeline's search query and filters.

Parameters

Name	In	Type	Req	Notes
`saved_search_id`	path	string	Yes	Saved-search id (returned by `POST /saved-searches`).

Responses

200 The updated saved-search object (same shape as the create response).

```
{
  "id": "<string>",
  "name": "<string>",
  "search_query": "<string>",
  "search_filters": { ... },
  "is_default": <boolean>,
  "created_at": "<timestamp>",
  "updated_at": "<timestamp>"
}
```

07 07 Entities & Review

8 endpoints

Review, vote on, split and remove the sources and datasets Ubiquity discovers.

#	Method	Path	Description
7.1	POST	`/locations/v1/entity-votes/toggle`	Like, dislike or flag an entity. Repeating the same vote type removes it.
7.2	GET	`/locations/v1/entity-votes/pipeline/{pipeline_id}`	List current votes for a pipeline.
7.3	GET	`/locations/api/v1/feeds/{feed_id}/entities`	List discovered entities (sources and datasets) for a feed, with search, sort and type filters.
7.4	DELETE	`/locations/api/v1/feeds/{feed_id}/entities/{entity_id}`	Remove an entity from a feed (hard delete).
7.5	GET	`/locations/api/v1/entities/{entity_id}/split-preview`	AI-cluster an entity's source URLs by topic, as a preview before splitting. Slow.
7.6	POST	`/locations/api/v1/entities/{entity_id}/split`	Commit a previewed split. Destructive and cross-feed.
7.7	GET	`/locations/api/v1/entities/{entity_id}`	Get a single entity's raw record.
7.8	GET	`/locations/api/v1/source-evaluation-logs/by-url`	Full AI evaluation reasoning for a discovered source, keyed by URL.

7.1 POST `/locations/v1/entity-votes/toggle`

Like, dislike or flag an entity. Repeating the same vote type removes it.

This is a true toggle: sending the same `vote_type` again removes the vote and returns `action: "removed"`.

`keycloak_user_id` here is the user UUID (the JWT `sub` claim), not the email. The chat endpoints use the opposite convention.

Request body `application/json`

Field	Type	Req	Notes
`entity_id`	integer		
`pipeline_id`	string		
`keycloak_user_id`	string		Keycloak user UUID (the JWT `sub` claim), not the email.
`vote_type`	string		One of: `like`, `dislike`, `flag`.

Example

```
{
  "entity_id": 0,
  "pipeline_id": "<pipeline_id>",
  "keycloak_user_id": "<keycloak_user_id>",
  "vote_type": "like"
}
```

Responses

200 Vote removed (the same `vote_type` was sent twice).

```
{
  "action": "removed",
  "vote": {...}
}
```

201 Vote created.

```
{
  "action": "created",
  "vote": {...}
}
```

7.2 GET `/locations/v1/entity-votes/pipeline/{pipeline_id}`

List current votes for a pipeline.

Parameters

Name	In	Type	Req	Notes
`pipeline_id`	path	string	Yes	Pipeline id.
`limit`	query	integer		Page size for list results.

Responses

200 Array of vote objects; `[]` if there are none.

7.3 GET `/locations/api/v1/feeds/{feed_id}/entities`

List discovered entities (sources and datasets) for a feed, with search, sort and type filters.

A CSV export of a feed is built on the client by combining this call (`limit=1000`, `offset=0`) with `GET /entity-votes/pipeline/{pipeline_id}?limit=500`. There is no dedicated export endpoint.

Parameters

Name	In	Type	Req	Notes
`feed_id`	path	integer	Yes	Numeric feed id (returned by `POST /feeds`).
`limit`	query	integer		Page size for list results.
`offset`	query	integer		Number of results to skip, for pagination.
`q`	query	string		
`order_by`	query	string		One of: `vote`, `name`, `type`, `created_at`.
`order`	query	string		One of: `asc`, `desc`.
`type`	query	string		`sources` is documented; `datasets` is not documented yet (expected to follow the same pattern). One of: `all`, `sources`, `datasets`. One of: `all`, `sources`, `datasets`.

Responses

200 Page of entities plus counts.

```
{
  "entities": [...],
  "total": <integer>,
  "dataset_count": <integer>,
  "source_count": <integer>
}
```

7.4 DELETE /locations/api/v1/feeds/{feed_id}/entities/{entity_id}

Remove an entity from a feed (hard delete).

Hard delete; it cannot be undone.

Parameters

Name	In	Type	Req	Notes
`feed_id`	path	integer	Yes	Numeric feed id (returned by `POST /feeds`).
`entity_id`	path	integer	Yes	Numeric entity id.

Responses

200 Entity removed.

```
{
  "message": "Entity removed from feed successfully"
}
```

7.5 GET /locations/api/v1/entities/{entity_id}/split-preview

AI-cluster an entity's source URLs by topic, as a preview before splitting. Slow.

Slow (several seconds) because it uses AI clustering; use a generous timeout.

`is_keeper` does not reliably land on the group that matches the entity's name, so inspect each preview.

A "No split needed" result can occasionally appear under rapid repeated calls; retry if you expect a split.

Parameters

Name	In	Type	Req	Notes
`entity_id`	path	integer	Yes	Numeric entity id.

Responses

200 Proposed grouping of the entity's sources.

```
{
  "entity_id": <integer>,
  "entity_name": "<string>",
  "total_sources": <integer>,
  "groups": [
    {
      "group_name": "<string>",
      "source_count": <integer>,
      "source_ids": [<integer>, ...],
      "urls": ["<string>", ...],
      "is_keeper": <boolean>
    }
  ],
  "discarded_urls": [...],
  "feeds_affected": [...]
}
```

7.6 POST /locations/api/v1/entities/{entity_id}/split

Commit a previewed split. Destructive and cross-feed.

Destructive and cross-feed. The original entity keeps the group marked `is_keeper` in the preview, and its name is not updated to match. Every other group becomes a new entity linked to every feed the original belonged to.

Parameters

Name	In	Type	Req	Notes
`entity_id`	path	integer	Yes	Numeric entity id.

Request body `application/json`

Field	Type	Req	Notes
`groups`	array of object		The groups returned by the preview, reduced to `group_name` and `source_ids`.
`groups[].group_name`	string		
`groups[].source_ids`	array of integer		

Example

```
{
  "groups": [
    {
      "group_name": "<group_name>",
      "source_ids": [
        0
      ]
    }
  ]
}
```

Responses

200 Split committed.

```
{
  "original_entity_id": <integer>,
  "original_entity_sources_kept": ...,
  "new_entities": [
    { "id": <integer>, "uuid": "<string>", "name": "<string>", "source_count": <integer> }
  ],
  "discarded_urls": [...],
  "feeds_linked": [...]
}
```

7.7 GET /locations/api/v1/entities/{entity_id}

Get a single entity's raw record.

Parameters

Name	In	Type	Req	Notes
`entity_id`	path	integer	Yes	Numeric entity id.

Responses

200 The entity's raw record: location, metadata (base64-encoded) and timestamps.

7.8 GET `/locations/api/v1/source-evaluation-logs/by-url`

Full AI evaluation reasoning for a discovered source, keyed by URL.

Parameters

Name	In	Type	Req	Notes
`url`	query	string	Yes	

Responses

200 The AI evaluation of the source.

```
{
  "id": ...,
  "credibility_score": ...,
  "local_focus": ...,
  "quality": ...,
  "evaluation_details": {
    "type": "<string>",
    "title": "<string>",
    "evaluation_notes": "<string>",
    ...
  }
}
```

08 08 Import Sources

3 endpoints

Bring your own URLs or behavioral datasets into a feed.

#	Method	Path	Description
8.1	POST	`/locations/api/v1/feeds/{feed_id}/sources/upload`	Import a list of source URLs (plain text or CSV flows). Async, returns 202.
8.2	POST	`/orchestrator/workflows/dataset-upload/trigger`	Import a behavioral dataset (multipart: jsonl, csv or xlsx, up to 25 MB).
8.3	GET	`/orchestrator/workflows/dataset-upload/{job_id}/status`	Poll a dataset upload job: pending, running, completed.

8.1 POST `/locations/api/v1/feeds/{feed_id}/sources/upload`

Import a list of source URLs (plain text or CSV flows). Async, returns 202.

Send either `urls` (plain text flow) or `urls_with_metadata` (CSV flow), not both. CSV parsing and preview happen on the client; this endpoint always receives the final URL list.

Returns 202 immediately. Progress polling for this `job_id` is not documented yet.

Parameters

Name	In	Type	Req	Notes
`feed_id`	path	integer	Yes	Numeric feed id (returned by `POST /feeds`).

Request body `application/json`

Field	Type	Req	Notes
`skip_validation`	boolean		
`location_uuid`	string		
`urls`	array of string		Plain Text sub-flow: array of raw URL strings.
`urls_with_metadata`	array of object		CSV Upload sub-flow: array of `{url, name?}` objects, mutually exclusive with `urls`.
`urls_with_metadata[].url`	string		
`urls_with_metadata[].name`	string		

Example

```
{
  "skip_validation": true,
  "location_uuid": "<location_uuid>",
  "urls": [
    "<urls>"
  ],
  "urls_with_metadata": [
    {
      "url": "<url>",
      "name": "<name>"
    }
  ]
}
```

Responses

202 Accepted for background processing.

```
{
  "job_id": "<string>",
  "feed_id": <integer>,
  "url_count": <integer>,
  "valid_urls": ...,
  "invalid_urls": ...,
  "status": "processing",
  "message": "<string>"
}
```

8.2 POST `/orchestrator/workflows/dataset-upload/trigger`

Import a behavioral dataset (multipart: jsonl, csv or xlsx, up to 25 MB).

Send as multipart form data with the form fields below plus the file (jsonl, csv or xlsx, up to 25 MB).

The name of the multipart part that carries the file is not documented yet.

Request body `multipart/form-data`

Field	Type	Req	Notes
`dataset_title`	string	Yes	Required.
`dataset_behavior`	string		Optional free-text tag, e.g. `consumer-spending`.
`dataset_unit`	string		Optional, e.g. `Percent of 2019`.
`dataset_source_url`	string (uri)		optional
`dataset_description`	string		Optional free text for analysts and Ubi.
`dataset_external_id`	string		Optional. Defaults to a value generated from the filename.

Example

```
dataset_title=<string>
dataset_behavior=consumer-spending
dataset_unit=Percent of 2019
dataset_source_url=https://example.com
dataset_description=<string>
dataset_external_id=<string>
<file part> (jsonl, csv or xlsx, max 25 MB; part name not documented yet)
```

Responses

200 Upload job created.

```
{
  "job_id": "<string>",
  "status": "pending",
  "progress": 0.0,
  "dataset_id": null,
  "series_names": [],
  "observations_stored": 0,
  "skipped_rows": 0,
  "warnings": [],
  "error": null,
  "created_at": "<timestamp>",
  "updated_at": "<timestamp>",
  "completed_at": null
}
```

8.3 GET `/orchestrator/workflows/dataset-upload/{job_id}/status`

Poll a dataset upload job: pending, running, completed.

Typical lifecycle: pending (progress 0.0), running (0.5), completed (1.0), reached within about 3 seconds for a trivial file.

Parameters

Name	In	Type	Req	Notes
`job_id`	path	string	Yes	Job id.

Responses

200 Same shape as the trigger response. On completion `dataset_id` is populated, `series_names` lists the file's non-date column headers, and `observations_stored`, `skipped_rows` and `warnings` are filled in.

```
{
  "job_id": "<string>",
  "status": "pending",
  "progress": 0.0,
  "dataset_id": null,
  "series_names": [],
  "observations_stored": 0,
  "skipped_rows": 0,
  "warnings": [],
  "error": null,
  "created_at": "<timestamp>",
  "updated_at": "<timestamp>",
  "completed_at": null
}
```

09 09 Discovery Jobs

7 endpoints

Launch, monitor and cancel agentic discovery work.

#	Method	Path	Description
9.1	POST	`/orchestrator/api/v1/jobs/create`	Launch a discovery job. The only call that actually starts agentic work; optional relative to pipeline creation.
9.2	GET	`/locations/api/v1/jobs/{job_id}`	Poll a job's status and progress.
9.3	PUT	`/locations/api/v1/jobs/{job_id}/status`	Cancel a job (only status=cancelled exercised).
9.4	GET	`/locations/api/v1/jobs`	List jobs across the account.
9.5	GET	`/locations/api/v1/jobs/count`	Total job count.
9.6	GET	`/locations/api/v1/jobs/count/status/{status}`	Job count by status.
9.7	GET	`/orchestrator/jobs/{job_id}/async-status`	Async sub-task progress for a job.

9.1 POST `/orchestrator/api/v1/jobs/create`

Launch a discovery job. The only call that actually starts agentic work; optional relative to pipeline creation.

Wait for a running job to finish before launching another for the same feed.

The request sends `job_type: "discovery"`; reading the job back reports `job_type: "curation"`. Both values are real.

Request body `application/json`

Field	Type	Req	Notes
`location_uuid`	string		
`job_type`	string		
`parameters`	object		
`parameters.search_query`	string		
`parameters.parsed_entities`	array of object		
`parameters.feed_id`	integer		
`parameters.pipeline_id`	string		
`parameters.vote_context`	object		
`parameters.vote_context.liked_entities`	array of object		Array of `{name, type}` objects used to steer discovery toward previously-liked entities.
`parameters.vote_context.liked_entities[].name`	string		
`parameters.vote_context.liked_entities[].type`	string		
`parameters.vote_context.disliked_entities`	array of object		
`max_queries`	integer		
`max_sources`	integer		
`min_sources`	integer		

Example

```
{
  "location_uuid": "<location_uuid>",
  "job_type": "discovery",
  "parameters": {
    "search_query": "<search_query>",
    "parsed_entities": [
      {}
    ],
    "feed_id": 0,
    "pipeline_id": "<pipeline_id>",
    "vote_context": {
      "liked_entities": [
        {
          "name": "<name>",
          "type": "<type>"
        }
      ],
      "disliked_entities": [
        {}
      ]
    }
  },
  "max_queries": 0,
  "max_sources": 0,
  "min_sources": 0
}
```

Responses

200 Job created and queued.

```
{
  "id": "<string>",
  "status": "pending",
  "progress": 0
}
```

9.2 GET `/locations/api/v1/jobs/{job_id}`

Poll a job's status and progress.

Parameters

Name	In	Type	Req	Notes
`job_id`	path	string	Yes	Job id.

Responses

200 Current job state.

```
{
  "id": "<string>",
  "status": "<pending | running | completed | failed | cancelled>",
  "progress": <0-100>,
  "job_type": "curation",
  "result_data": ...,
  "pipeline_id": "<string>"
}
```

9.3 PUT /locations/api/v1/jobs/{job_id}/status

Cancel a job (only status=cancelled exercised).

Only `status: "cancelled"` has been exercised.

Parameters

Name	In	Type	Req	Notes
job_id	path	string	Yes	Job id.

Request body `application/json`

Field	Type	Req	Notes
status	string		

Example

```
{
  "status": "cancelled"
}
```

Responses

200 The job, echoed with `status` set to `cancelled` and `completed_at` set.

9.4 GET /locations/api/v1/jobs

List jobs across the account.

Parameters

Name	In	Type	Req	Notes
limit	query	integer		Page size for list results.
offset	query	integer		Number of results to skip, for pagination.

Responses

200 Array of job objects (same shape as `GET /jobs/{job_id}`).

9.5 GET /locations/api/v1/jobs/count

Total job count.

Responses

200 Total number of jobs.

```
{
  "count": <integer>
}
```

9.6 GET `/locations/api/v1/jobs/count/status/{status}`

Job count by status.

Parameters

Name	In	Type	Req	Notes
`status`	path	string	Yes	Job status to count. One of: `pending`, `running`, `completed`, `failed`, `cancelled`. One of: `pending`, `running`, `completed`, `failed`, `cancelled`.

Responses

200 Number of jobs in that status.

```
{
  "count": <integer>
}
```

9.7 GET `/orchestrator/jobs/{job_id}/async-status`

Async sub-task progress for a job.

Parameters

Name	In	Type	Req	Notes
`job_id`	path	string	Yes	Job id.

Responses

200 Progress of the job's asynchronous sub-tasks.

```
{
  "job_id": "<string>",
  "pending_count": <integer>,
  "processed_count": <integer>,
  "failed_count": <integer>,
  "total_count": <integer>,
  "all_complete": <boolean>,
  "has_async_work": <boolean>,
  "snapshots": [...]
}
```

10 10 Understand

9 endpoints

Per-feed chat with Ubi, plus artifact and metric data for the Data Browser.

#	Method	Path	Description
10.1	GET	`/understand/chat/sessions`	List Ubi chat sessions for a user and feed.
10.2	POST	`/understand/chat/sessions`	Create a chat session (created lazily on first message).
10.3	GET	`/understand/chat/sessions/{session_id}`	Get one session's metadata.
10.4	DELETE	`/understand/chat/sessions/{session_id}`	Delete a chat session.
10.5	GET	`/understand/chat/sessions/{session_id}/messages`	Read a session's message history.
10.6	POST	`/understand/chat/sessions/{session_id}/messages`	Send a message to Ubi. Synchronous: the reply is in the same response. Required field: message.
10.7	GET	`/understand/artifacts/count/{feed_id}`	Count ingested artifacts for a feed. Preferred fast ingestion-progress check.

#	Method	Path	Description
10.8	POST	`/understand/artifacts/preview`	Fetch artifacts for the Data Browser. Capped at 1000 items regardless of limit.
10.9	POST	`/understand/metrics/preview`	Custom metric time series and per-artifact scores. Slow and uncapped; always scope by date range.

10.1 GET `/understand/chat/sessions`

List Ubi chat sessions for a user and feed.

`keycloak_user_id` is the account email here, the opposite of the entity-votes call.

Parameters

Name	In	Type	Req	Notes
`keycloak_user_id`	query	string	Yes	The account email, not the Keycloak UUID.
`feed_id`	query	integer	Yes	Numeric feed id (returned by `POST /feeds`).
`limit`	query	integer		Page size for list results.
`offset`	query	integer		Number of results to skip, for pagination.

Responses

200 `[]`, or an array of session objects. Each session's `filter_state` is derived automatically from its first message.

```
{
  "id": "<string>",
  "feed_id": <integer>,
  "title": "<string>",
  "filter_state": {
    "topics": [...], "entities": [...], "dateRange": {...},
    "siteTypes": [...], "metricName": "<string>",
    "textSearch": "<string>", "entityTypes": [...]
  }
}
```

10.2 POST `/understand/chat/sessions`

Create a chat session (created lazily on first message).

The session is created lazily, on the first message, rather than when a dedicated "New Chat" action is clicked.

Responses

200 Session created.

```
{
  "id": "<string>",
  "feed_id": <integer>,
  "title": "<string>",
  "filter_state": {
    "topics": [...], "entities": [...], "dateRange": {...},
    "siteTypes": [...], "metricName": "<string>",
    "textSearch": "<string>", "entityTypes": [...]
  }
}
```

10.3 GET /understand/chat/sessions/{session_id}

Get one session's metadata.

Parameters

Name	In	Type	Req	Notes
session_id	path	string	Yes	Chat session id.
keycloak_user_id	query	string	Yes	

Responses

200 The session.

```
{
  "id": "<string>",
  "feed_id": <integer>,
  "title": "<string>",
  "filter_state": {
    "topics": [...], "entities": [...], "dateRange": {...},
    "siteTypes": [...], "metricName": "<string>",
    "textSearch": "<string>", "entityTypes": [...]
  }
}
```

10.4 DELETE /understand/chat/sessions/{session_id}

Delete a chat session.

Parameters

Name	In	Type	Req	Notes
session_id	path	string	Yes	Chat session id.

Responses

200 Session deleted.

```
{
  "status": "deleted",
  "session_id": "<string>"
}
```

10.5 GET /understand/chat/sessions/{session_id}/messages

Read a session's message history.

Parameters

Name	In	Type	Req	Notes
session_id	path	string	Yes	Chat session id.
keycloak_user_id	query	string	Yes	
limit	query	integer		Page size for list results.
offset	query	integer		Number of results to skip, for pagination.

Responses

200 Array of messages, oldest to newest as stored.

```
{
  "id": "<string>",
  "session_id": "<string>",
  "role": "<user | assistant>",
  "content": "<string>",
  "citations": [ { "id": ..., "url": "<string>", "score": ..., "title": "<string>" } ],
  "chart": { "type": "<string>", "title": "<string>", "config": {...} },
  "created_at": "<timestamp>"
}
```

10.6 POST `/understand/chat/sessions/{session_id}/messages`

Send a message to Ubi. Synchronous: the reply is in the same response. Required field: message.

Synchronous: the reply comes back in this response and no polling is needed. A fresh session is created implicitly on the first call.

A trivial question takes about 3 to 4 seconds. Complex, feed-wide questions can take several minutes, so set a long client timeout.

This is a different chat system from the refine loop on `POST /search/unified`: this one is the standing per-feed Understand chat.

Request body `application/json`

Field	Type	Req	Notes
`message`	string	Yes	Required. The field is `message`, not `content`; an empty body returns 422.

Example

```
{
  "message": "<message>"
}
```

Responses

200 Ubi's reply, returned inline.

```
{
  "user_message": {
    "id": "<string>", "session_id": "<string>", "role": "user",
    "content": "<string>", "created_at": "<timestamp>"
  },
  "assistant_message": {
    "id": "<string>", "role": "assistant", "content": "<string>",
    "reasoning_summary": "<string>", "citations": [...], "chart": {...},
    "created_at": "<timestamp>"
  },
  "citations": [...],
  "artifacts_count": <integer>,
  "chart": {...},
  "filter_update": {
    "dateRange": {...}, "siteTypes": [...], "topics": [...],
    "entities": [...], "textSearch": "<string>", "reason": "<string>"
  }
}
```

10.7 GET /understand/artifacts/count/{feed_id}

Count ingested artifacts for a feed. Preferred fast ingestion-progress check.

A fast, inexpensive way to track ingestion progress; cheaper than reading `total_count` from `POST /artifacts/preview`.

Parameters

Name	In	Type	Req	Notes
feed_id	path	integer	Yes	Numeric feed id (returned by `POST /feeds`).

Responses

200 Number of artifacts ingested for the feed.

```
{
  "count": <integer>,
  "feed_id": <integer>
}
```

10.8 POST /understand/artifacts/preview

Fetch artifacts for the Data Browser. Capped at 1000 items regardless of limit.

`limit` is accepted but ignored: at most 1000 items are ever returned.

Request body `application/json`

Field	Type	Req	Notes
feed_id	integer		
filter_state	object		
filter_state.dateRange	object		
filter_state.dateRange.start	string (date)		
filter_state.dateRange.end	string (date)		
filter_state.topics	array		May be null. May be null.
filter_state.entities	array		May be null. May be null.
filter_state.siteTypes	array of string		
filter_state.metricName	string		May be null. May be null.
filter_state.textSearch	string		May be null. May be null.
filter_state.entityTypes	array		May be null. May be null.
limit	integer		Accepted but ignored; at most 1000 items are returned.

Example

```
{
  "feed_id": 0,
  "filter_state": {
    "dateRange": {
      "start": "YYYY-MM-DD",
      "end": "YYYY-MM-DD"
    },
    "topics": [],
    "entities": [],
    "siteTypes": [
      "<siteTypes>"
    ],
    "metricName": "<metricName>",
    "textSearch": "<textSearch>",
    "entityTypes": []
  },
  "limit": 0
}
```

Responses

200 Up to 1000 matching artifacts.

```
{
  "artifacts": [
    {
      "id": "<string>",
      "headline": "<string>",
      "url": "<string>",
      "site_type": "<string>",
      "published_at": "<timestamp>",
      "content_sample": "<string>",
      "score": ...
    }
  ],
  "total_count": <integer>
}
```

500 Returned as "Collection <email> doesn't exist" before any artifacts have been ingested for the account. This is expected at that stage and is not a real failure.

10.9 POST `/understand/metrics/preview`

Custom metric time series and per-artifact scores. Slow and uncapped; always scope by date range.

No result cap, so it is slow (roughly 45 to 150+ seconds). Always scope the request with `filter_state.dateRange` and use a long timeout.

Request body `application/json`

Field	Type	Req	Notes
`feed_id`	integer		
`filter_state`	object		
`filter_state.metricName`	string		
`filter_state.dateRange`	object		
`filter_state.dateRange.start`	string (date)		
`filter_state.dateRange.end`	string (date)		

Example

```
{
  "feed_id": 0,
  "filter_state": {
    "metricName": "<metricName>",
    "dateRange": {
      "start": "YYYY-MM-DD",
      "end": "YYYY-MM-DD"
    }
  }
}
```

Responses

200 Time series plus per-artifact scores.

```
{
  "data": [ { "artifact_uuid": "<string>", "score": ..., "list": ... } ],
  "series": [ { "bucket": "<string>", "score": ... } ],
  "total_count": <integer>,
  "filter_applied": {...}
}
```